

W4995 Applied Machine Learning

Trees, Forests and Boosting

02/20/17


Andreas Müller

Why trees?

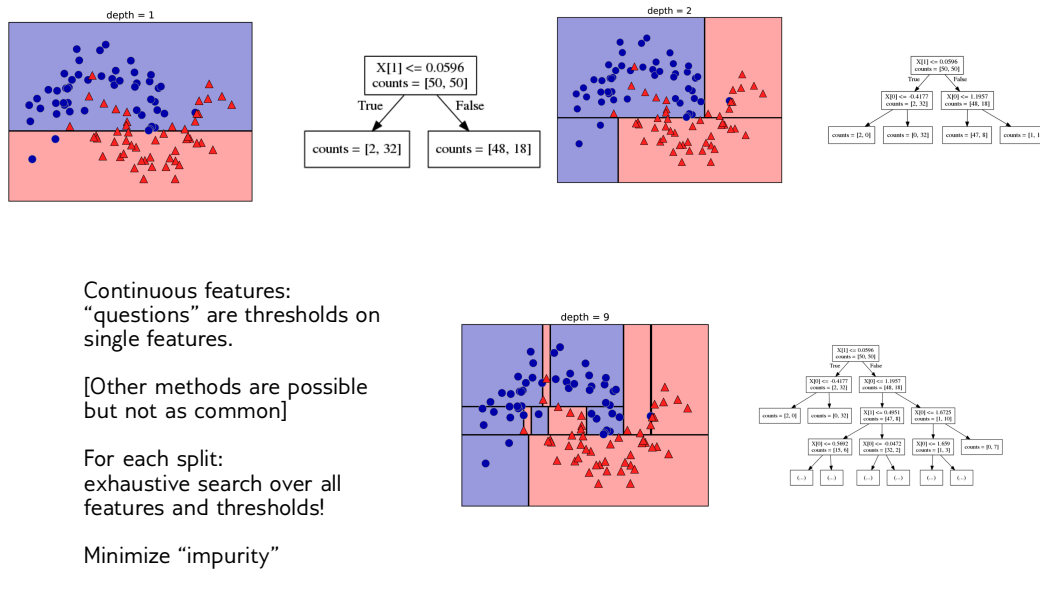
- Very powerful modeling method – non-linear!
- Doesn't care about scaling!
- “Interpretable”
- Basis of very powerful models!

Decision Trees for Classification

Idea: series of binary questions



Building trees



The second family of models I want to talk about is decision-tree based models.

A decision tree is basically the same as a sequence of if-else branches, that ultimately lead to a decision. The point is that the questions that are asked are learned, though.

A decision tree is build recursively by asking a series of questions of the form “is feature “i” greater than threshold t”. In each iteration, the question is chosen that yields the most information about the target variable. Then, the data is split according to this question, and we start again. This yields a hierarchical partitioning of the data, where each section of the partitioning becomes more and more “pure”, that is their content becomes more and more the same.

After you build the tree, you can make a prediction by checking which part of the partition a new point lies in and assigning the mean of the datapoints in this part.

Criteria (for classification)

- Gini index:

$$H_{\text{gini}}(X_m) = \sum_{k \in \mathcal{Y}} p_{mk}(1 - p_{mk})$$

- Cross-entropy:



$$H_{\text{CE}}(X_m) = - \sum_{k \in \mathcal{Y}} p_{mk} \log(p_{mk})$$

X_m observations in node m

\mathcal{Y} classes

$p_{m\cdot}$ distribution over classes in node m

Prediction



Traverse tree based on feature tests
Predict most common class in leaf

Regression trees

- Impurity Criteria:
Mean Squared Error
Mean Absolute Error
- Prediction:
Predict mean.
- Without regularization / pruning:
Each leaf often contains a single point to be “pure”

Visualizing trees with sklearn

```
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
```

```
X_train, X_test, y_train, y_test = train_test_split(cancer.data, cancer.target, stratify=cancer.target, random_state=0)
```

```
# tree visualization
```

```
from sklearn.tree import DecisionTreeClassifier, export_graphviz
tree = DecisionTreeClassifier(max_depth=2)
tree.fit(X_train, y_train)
```

```
tree_dot = export_graphviz(tree, out_file=None, feature_names=cancer.feature_names)
```

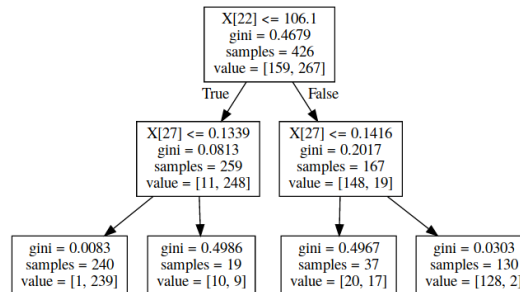
```
print(tree_dot)
```

```
digraph Tree {
  node [shape=box] ;
  0 [label="X[22] <= 106.1\ngini = 0.4679\nsamples = 426\nvalue = [159, 267]" ;
  1 [label="X[27] <= 0.1339\ngini = 0.0813\nsamples = 259\nvalue = [11, 248]" ;
  0 -> 1 [labeldistance=2.5, labelangle=45, headlabel="True"] ;
  2 [label="gini = 0.0083\nsamples = 240\nvalue = [1, 239]" ;
  1 -> 2 ;
  3 [label="gini = 0.4986\nsamples = 19\nvalue = [10, 9]" ;
  1 -> 3 ;
  4 [label="X[27] <= 0.1416\ngini = 0.2017\nsamples = 167\nvalue = [148, 19]" ;
  0 -> 4 [labeldistance=2.5, labelangle=-45, headlabel="False"] ;
  5 [label="gini = 0.4967\nsamples = 37\nvalue = [20, 17]" ;
  4 -> 5 ;
  6 [label="gini = 0.0303\nsamples = 130\nvalue = [128, 2]" ;
  4 -> 6 ;
}
```

Showing dot files in Jupyter

- First install graphviz C library:
conda install graphviz
- Then install graphviz python library:
pip install graphviz

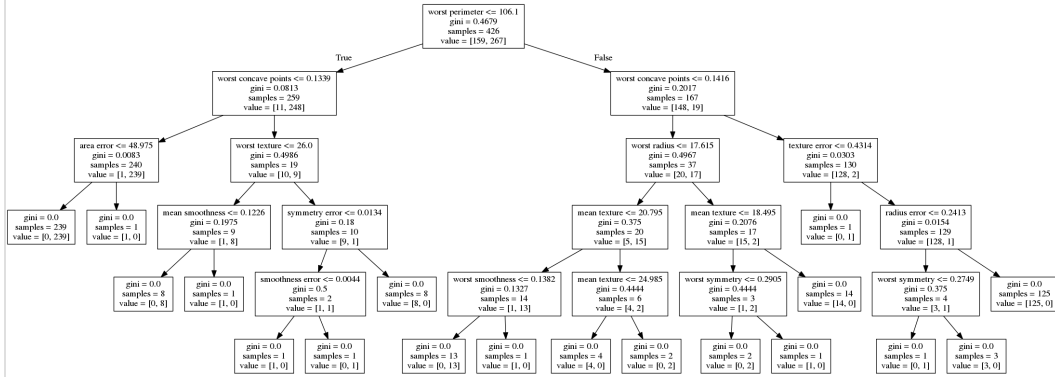
```
import graphviz
graphviz.Source(tree_dot)
```



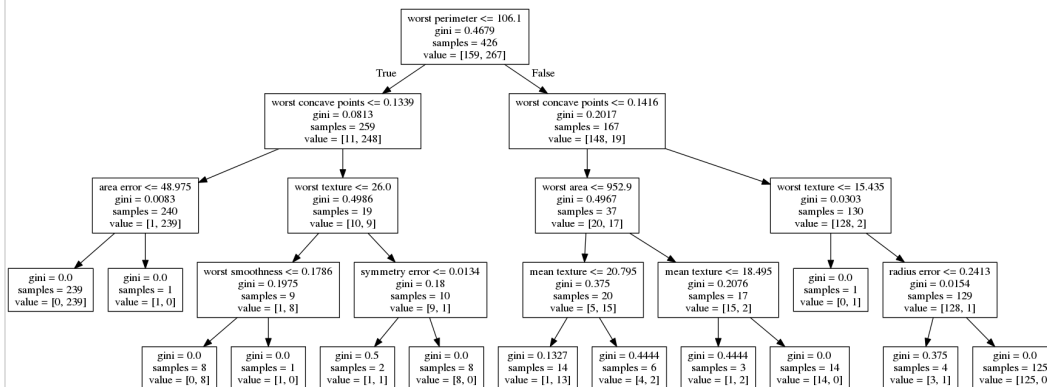
Parameter tuning

- Pre-pruning and post-pruning (not in sklearn yet)
- Limit tree size (pick one):
 - `max_depth`
 - `max_leaf_nodes`
 - `min_samples_split`
 - (and more)

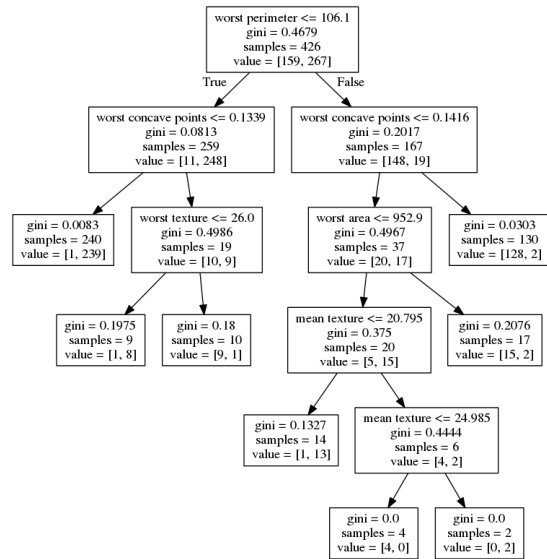
No pruning



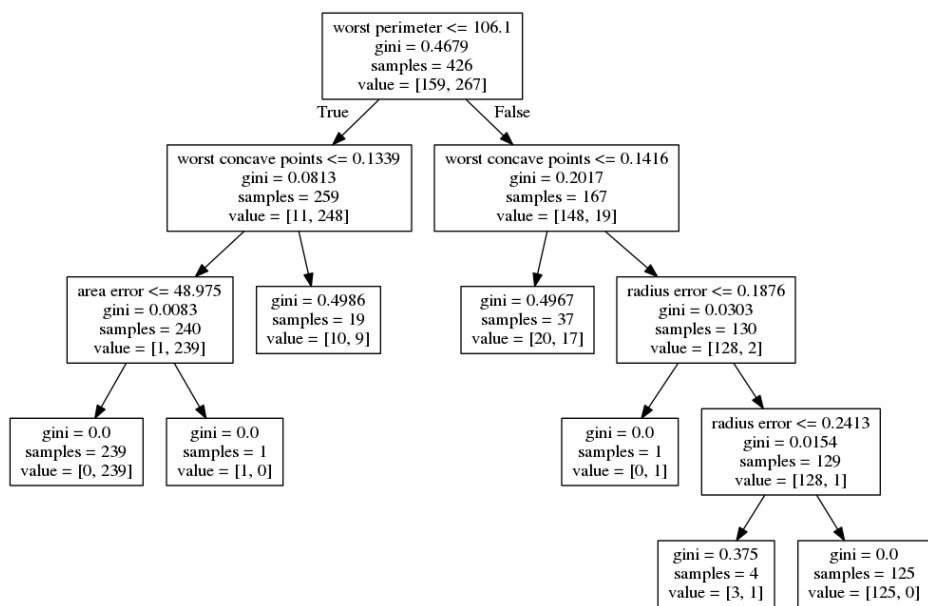
max_depth=4




max_leaf_nodes=8




min_samples_split=50





```
from sklearn.model_selection import GridSearchCV
param_grid = {'max_depth': range(1, 7)}
grid = GridSearchCV(DecisionTreeClassifier(random_state=0), param_grid=param_grid, cv=10)
grid.fit(X_train, y_train)
```





```
from sklearn.model_selection import GridSearchCV
param_grid = {'max_leaf_nodes':range(2, 20)}
grid = GridSearchCV(DecisionTreeClassifier(random_state=0), param_grid=param_grid, cv=10)
grid.fit(X_train, y_train)
```







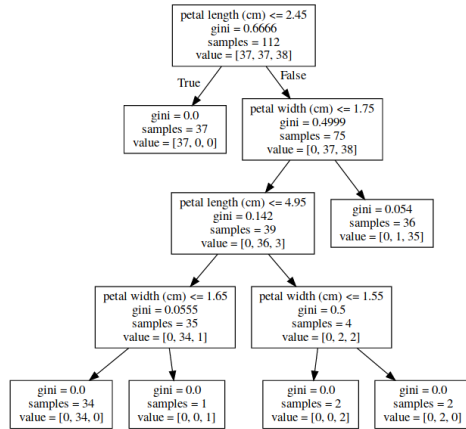
Extrapolation



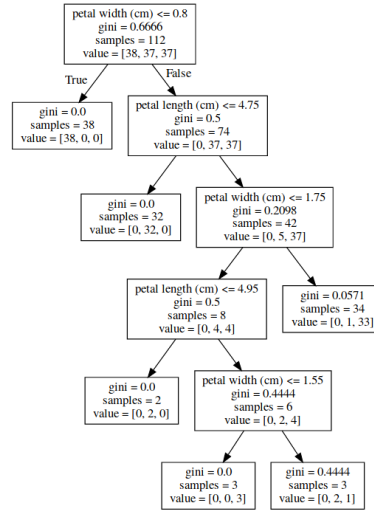
Would be the same for nearest neighbor regression!

Instability

```
iris = load_iris()
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, stratify=iris.target, random_state=0)
tree = DecisionTreeClassifier(max_leaf_nodes=6).fit(X_train, y_train)
tree_dot = export_graphviz(tree, out_file=None, feature_names=iris.feature_names)
graphviz.Source(tree_dot)
```

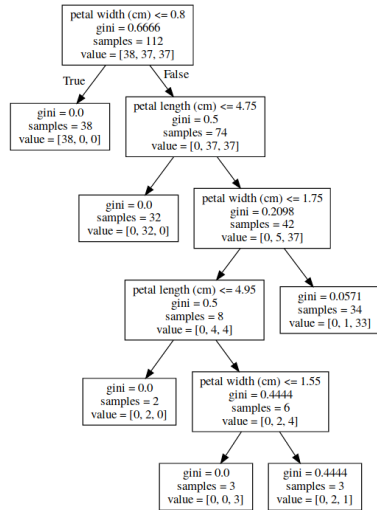


```
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, stratify=iris.target, random_state=1)
tree = DecisionTreeClassifier(max_leaf_nodes=6).fit(X_train, y_train)
tree_dot = export_graphviz(tree, out_file=None, feature_names=iris.feature_names)
graphviz.Source(tree_dot)
```




Feature importance

```
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, stratify=iris.target, random_state=1)
tree = DecisionTreeClassifier(max_leaf_nodes=6).fit(X_train, y_train)
tree_dot = export_graphviz(tree, out_file=None, feature_names=iris.feature_names)
graphviz.Source(tree_dot)
```



tree.feature_importances_

array([0. , 0. , 0.414, 0.586])



Unstable tree → unstable feature importances.

Might take one or multiple from a group of correlated features.

Categorical Data

- Can split on categorical data directly
- Intuitive way to split: split in two subsets
- 2^{n_values} many possibilities
- Possible to do in linear time exactly for gini index and binary classification.
- Heuristics done in practice for multi-class.
- Not in sklearn release version :(

Predicting probabilities

- Fraction of class in leaf.
- Without pruning: Always 100% certain!
- Even with pruning might be too certain.

Conditional Inference Trees


- Select “best” split with correcting for multiple-hypothesis testing.
- More “fair” to categorical variables.
- Only in R so far (party)

Relation to Nearest Neighbors

- Predict average of neighbors – either by k , by epsilon ball or by leaf.
- Trees are much faster to predict.
- Both can't extrapolate

Different Splitting Methods

- Could use anything as split candidate!
- Linear models used if extrapolation is needed.
- Computer vision: pixel comparisons
- Kinect (first generation): depth comparison



(taken from Shotton et. al. Real-Time Human Pose Recognition ..)

Ensemble Models




Poor man's ensembles

- Build different models
- Average the result
- Owen Zhang (long time kaggle 1st): build XGBoosting models with different random seeds.
- More models are better – if they are not correlated.
- Also works with neural networks
- You can average any models as long as they provide calibrated (“good”) probabilities.
- Scikit-learn: VotingClassifier
hard and soft voting

VotingClassifier


```
voting = VotingClassifier([('logreg', LogisticRegression(C=100)),
                           ('tree', DecisionTreeClassifier(max_depth=3, random_state=0))],
                           voting='soft')
voting.fit(X_train, y_train)
lr, tree = voting.estimators_
print("{:.2f} " * 3).format(voting.score(X_test, y_test),
                             lr.score(X_test, y_test), tree.score(X_test, y_test)))
```

0.88 0.84 0.80




Bagging (Bootstrap AGGregation)

- Generic way to build “slightly different” models
- Draw bootstrap samples from dataset
(as many as there are in the dataset, with repetition)
- Implemented in BaggingClassifier, BaggingRegressor



Bias and Variance




<http://scott.fortmann-roe.com/docs/BiasVariance.html>

Bias and Variance in Ensembles

- Breiman showed that generalization depends on strength of the individual classifiers and (inversely) on their correlation
- Uncorrelating them might help, even at the expense of strength

Random Forests

- Smarter bagging for trees!




So decision trees are a great idea, but unfortunately they don't work that well in practice. However, there is a modification of the algorithm that works very well, called random forest.

The idea behind random forest is that we build many decision trees, but we inject some randomness into each tree, so that they are all different.


Then, to make a prediction, we look at the prediction of all the decision trees and take the average.

Randomize in two ways

- For each **tree**:
Pick bootstrap sample of data



- For each **split**:
Pick random sample of features



- More tree are always better

Tuning Random Forests

- Main parameter: max_features
 - around $\sqrt{n_features}$ for classification
 - Around $n_features$ for regression
- n_estimators > 100
- Prepruning might help, definitely helps with model size!
- max_depth, max_leaf_nodes, min_samples_split again

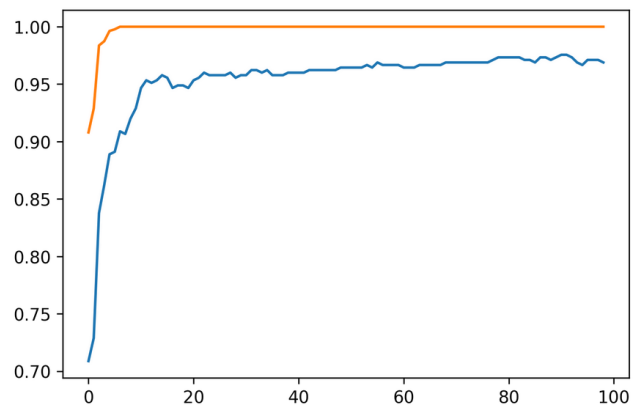
Extremely Randomized Trees

- More randomness!
- Randomly draw threshold for each feature!
- Doesn't use bootstrap
- Faster because no sorting / searching
- Can have smoother boundaries

Warm-Starts

```
train_scores = []
test_scores = []


rf = RandomForestClassifier(warm_start=True)
for n_estimators in range(1, 100, 5):
    rf.n_estimators = n_estimators
    rf.fit(X_train, y_train)
    train_scores.append(rf.score(X_train, y_train))
    test_scores.append(rf.score(X_test, y_test))
```



Out-of-bag estimates

- Each tree only uses ~66% of data
- Can evaluate it on the rest!
- Make predictions for out-of-bag, average, score.
- Each prediction is an average over different subset of trees

```
feature_range = range(1, 64, 5)
for max_features in feature_range:
    rf = RandomForestClassifier(max_features=max_features, oob_score=True, n_estimators=200, random_state=0)
    rf.fit(X_train, y_train)
    train_scores.append(rf.score(X_train, y_train))
    test_scores.append(rf.score(X_test, y_test))
    oob_scores.append(rf.oob_score_)
```




Variable Importance

```
X_train, X_test, y_train, y_test = train_test_split(  
    iris.data, iris.target, stratify=iris.target, random_state=1)  
rf = RandomForestClassifier(n_estimators=100).fit(X_train, y_train)
```


```
rf.feature_importances_  
array([ 0.101,  0.034,  0.437,  0.428])
```

```
plt.barh(range(4), rf.feature_importances_)  
plt.yticks(range(4), iris.feature_names);
```




Gradient Boosting

Gradient Boosting Algorithm



- Iteratively add regression trees to model
- Use log loss for classification
- Discount update by learning rate

GradientBoostingClassifier(max_depth=2)



Gradient Boosting

- Many shallow trees
- `learning_rate` \leftrightarrow `n_estimators`
- Slower to train than RF (serial), but much faster to predict
- Small model size
- Uses one-vs-rest for multi-class!

Tuning Gradient Boosting

- Pick `n_estimators`, tune learning rate
- Can also tune `max_features`
- Typically strong pruning via `max_depth`

Partial Dependence Plots


- Marginal dependence of prediction on one or two features

```
boston = load_boston()
X_train, X_test, y_train, y_test = train_test_split(
    boston.data, boston.target, random_state=0)


gbrt = GradientBoostingRegressor().fit(X_train, y_train)
gbrt.score(X_test, y_test)
```

0.81962024379538989

```
from sklearn.ensemble.partial_dependence import plot_partial_dependence
fig, axs = plot_partial_dependence(gbrt, X_train, np.argsort(gbrt.feature_importances_)[-6:],
    feature_names=boston.feature_names,
    n_jobs=5, grid_resolution=50)
plt.tight_layout()
```

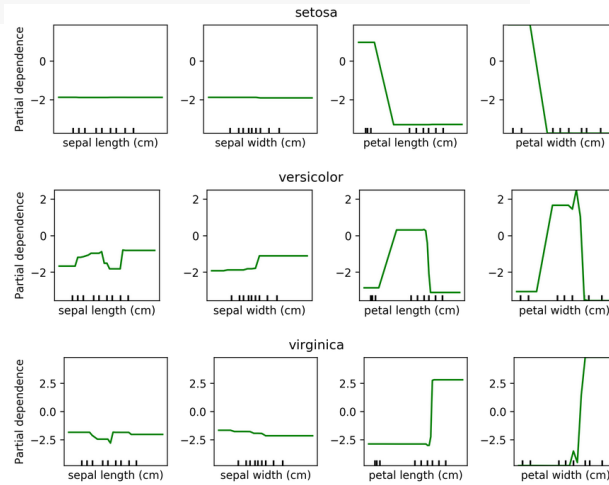


```
fig, axes = plot_partial_dependence(gbrt, X_train, [np.argsort(gbrt.feature_importances_)[-2:]],  
                                   feature_names=boston.feature_names,  
                                   n_jobs=3, grid_resolution=50)
```



Partial Dependence for Classification

```
from sklearn.ensemble.partial_dependence import plot_partial_dependence
for i in range(3):
    fig, axs = plot_partial_dependence(gbrt, X_train, range(4), n_cols=4,
                                       feature_names=iris.feature_names, grid_resolution=50, label=i,
                                       figsize=(8, 2))
    fig.suptitle(iris.target_names[i])
    for ax in axs: ax.set_xticks(())
plt.tight_layout()
```



XGBoost

- Efficient implementation of gradient boosting
- Improvements on original algorithm
- <https://arxiv.org/abs/1603.02754>
- Adds l1 and l2 penalty on leaf-weights
- Fast approximate split finding
- Can pip-install
- Scikit-learn compatible interface

Boosting in General

- “Meta-algorithm” to create strong learners from weak learners.
- AdaBoost, GentleBoost, ...
- Trees or stumps work best
- Gradient Boosting often the best of the bunch
- Many specialized algorithms (ranking etc)

When to use tree-based models

- Model non-linear relationships
- Single tree: very interpretable (if small)
- Random forests very robust, good benchmark
- Gradient boosting often best performance with careful tuning
- Doesn't care about scaling, no need for feature engineering!

More ensembles: Stacking


Poor man's Stacking

- Build multiple models
- Train model on probabilities / scores produced

```
from sklearn.neighbors import KNeighborsClassifier
X, y = make_moons(noise=.2, random_state=18)
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, random_state=0)

voting = VotingClassifier([('logreg', LogisticRegression(C=100)),
                           ('tree', DecisionTreeClassifier(max_depth=3, random_state=0)),
                           ('knn', KNeighborsClassifier(n_neighbors=3))
                          ],
                          voting='soft')
voting.fit(X_train, y_train)
lr, tree, knn = voting.estimators_
print("{:.2f} " * 4).format(voting.score(X_test, y_test),
                           lr.score(X_test, y_test), tree.score(X_test, y_test),
                           knn.score(X_test, y_test))

0.88 0.84 0.80 1.00
```





Poor man's Stacking

```
from sklearn.preprocessing import FunctionTransformer
# we need to reshape the result from votingclassifier.transform because
# of some annoyance in sklearn. We then keep only the probabilities of the positive classes!
reshaper = FunctionTransformer(lambda X: np.rollaxis(X, 1).reshape(-1, 6)[: , 1::2], validate=False)
stacking = make_pipeline(voting, reshaper,
                        LogisticRegression(C=100))
stacking.fit(X_train, y_train)
stacking.score(X_train, y_train)
0.9866666666666669

stacking.score(X_test, y_test)
0.95999999999999996

stacking.named_steps['logisticregression'].coef_
array([[ -2.625,   6.261,   9.501]])
```


Problem: Overfitting!



Stacking

- Use cross-validation (even LOO!) to produce probability estimates on training set.
- Train second step estimator on held-out estimates
- No overfitting of second step!
- For testing: as usual

Hold-out estimates of probabilities



- Split 1 produces probabilities for Fold 1, split2 for Fold 2 etc.
- Get a probability estimate for each data point!
- Unbiased estimates (like on the test set) for the whole training set!
- Without it: The best estimator is the one that memorized the training set.

Stacking with Scikit-learn


```
from sklearn.model_selection import cross_val_predict
first_stage = make_pipeline(voting, reshaper)
transform_cv = cross_val_predict(first_stage, X_train, y_train, cv=10, method="transform")
```

```
second_stage = LogisticRegression(C=100).fit(transform_cv, y_train)
print(second_stage.coef_)
[[ 2.09 -1.424  7.93 ]]
```


```
second_stage.score(transform_cv, y_train)
0.95999999999999996
```

```
second_stage.score(first_stage.transform(X_test), y_test)
1.0
```

Hold-out stacking



Naive stacking



Summary